

# Medical\_Insurance\_Prediction

December 31, 2024

Anthony Adungo Machine Learning Project.

## 1 Project Description

This project focuses on predicting medical insurance premiums using a publicly available dataset. The dataset includes key attributes such as demographics, health-related factors, and personal habits, which influence the cost of insurance premiums. The primary objective is to develop machine-learning models capable of accurately forecasting premium amounts. This involves identifying the most significant factors affecting premiums and providing actionable insights to refine pricing strategies, enabling insurers to optimize their offerings and enhance customer satisfaction.

## 2 Key Objectives

Data Exploration and Preprocessing.

Perform exploratory data analysis (EDA) to understand the dataset's structure. Clean the data by handling missing values, addressing outliers, and encoding categorical variables where necessary.

Feature Engineering.

Develop and refine features to enhance model performance by identifying relationships between demographic, behavioral, and other relevant factors.

Model Development

Utilize machine learning algorithms, including Multiple Linear Regression, Random Forest Regressor, and XGBoost Regressor, to predict medical insurance premiums. Compare the performance of these models to determine the best approach for accurate predictions.

Model Evaluation

Evaluate each model's performance using metrics such as  $R^2$ , Mean Absolute Error (MAE), and Root Mean Squared Error (RMSE). Ensure the models are robust and capable of generalizing well to unseen data.

Insights and Recommendations

Analyze model outputs to identify the most influential factors affecting insurance premiums. Provide actionable insights for insurers to optimize premium pricing strategies and enhance customer satisfaction.

### 3 Tools and Technologies

Programming: Python (Pandas, NumPy, Scikit-learn, XGBoost) Environment: Jupyter Notebooks  
Visualization: Matplotlib, Seaborn Machine Learning: Multiple Linear Regression, Random Forest Regressor, XGBoost Regressor

### 4 Part 1: Data Pre-processing

I) Importing the dataset and exploring its properties.

```
[95]: #Importing all the necessary libraries.  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
%matplotlib inline  
import seaborn as sns  
import warnings  
warnings.filterwarnings('ignore')  
from scipy import stats
```

```
[96]: df = pd.read_csv('insurance.csv')  
df.sample(5)
```

```
[96]:
```

	age	sex	bmi	children	smoker	region	charges
1245	28	male	24.300	5	no	southwest	5615.36900
908	63	male	39.800	3	no	southwest	15170.06900
319	32	male	37.335	1	no	northeast	4667.60765
462	62	female	38.095	2	no	northeast	15230.32405
729	41	female	36.080	1	no	southeast	6781.35420

```
[97]: df.shape
```

```
[97]: (1338, 7)
```

```
[98]: df.columns
```

```
[98]: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'],  
dtype='object')
```

```
[99]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1338 entries, 0 to 1337  
Data columns (total 7 columns):  
#   Column      Non-Null Count  Dtype  
---  -  
0   age         1338 non-null   int64  
1   sex         1338 non-null   object  
2   bmi         1338 non-null   float64
```

```

3  children  1338 non-null  int64
4  smoker    1338 non-null  object
5  region    1338 non-null  object
6  charges   1338 non-null  float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB

```

```
[100]: #Finding the statistical summary of the dataset.
df.describe().transpose()
```

```
[100]:
```

	count	mean	std	min	25%	50% \
age	1338.0	39.207025	14.049960	18.0000	27.00000	39.000
bmi	1338.0	30.663397	6.098187	15.9600	26.29625	30.400
children	1338.0	1.094918	1.205493	0.0000	0.00000	1.000
charges	1338.0	13270.422265	12110.011237	1121.8739	4740.28715	9382.033

	75%	max
age	51.000000	64.00000
bmi	34.693750	53.13000
children	2.000000	5.00000
charges	16639.912515	63770.42801

```
[101]: print(df.dtypes)
```

```

age          int64
sex          object
bmi         float64
children     int64
smoker       object
region       object
charges     float64
dtype: object

```

```
[102]: #Finding the categorical variables.
df.select_dtypes(include='object').columns
```

```
[102]: Index(['sex', 'smoker', 'region'], dtype='object')
```

```
[103]: len(df.select_dtypes(include='object').columns)
```

```
[103]: 3
```

```
[104]: df.select_dtypes(include=['float64','int64']).columns
```

```
[104]: Index(['age', 'bmi', 'children', 'charges'], dtype='object')
```

```
[105]: len(df.select_dtypes(include=['float64','int64']).columns)
```

```
[105]: 4
```

## II) Dealing with missing values.

```
[106]: missing_data = df.isnull()
print(missing_data.head())
for column in missing_data.columns.values.tolist():
    print(column)
    print(missing_data[column].value_counts())
    print(" ")
```

	age	sex	bmi	children	smoker	region	charges
0	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False

age

age

False 1338

Name: count, dtype: int64

sex

sex

False 1338

Name: count, dtype: int64

bmi

bmi

False 1338

Name: count, dtype: int64

children

children

False 1338

Name: count, dtype: int64

smoker

smoker

False 1338

Name: count, dtype: int64

region

region

False 1338

Name: count, dtype: int64

charges

charges

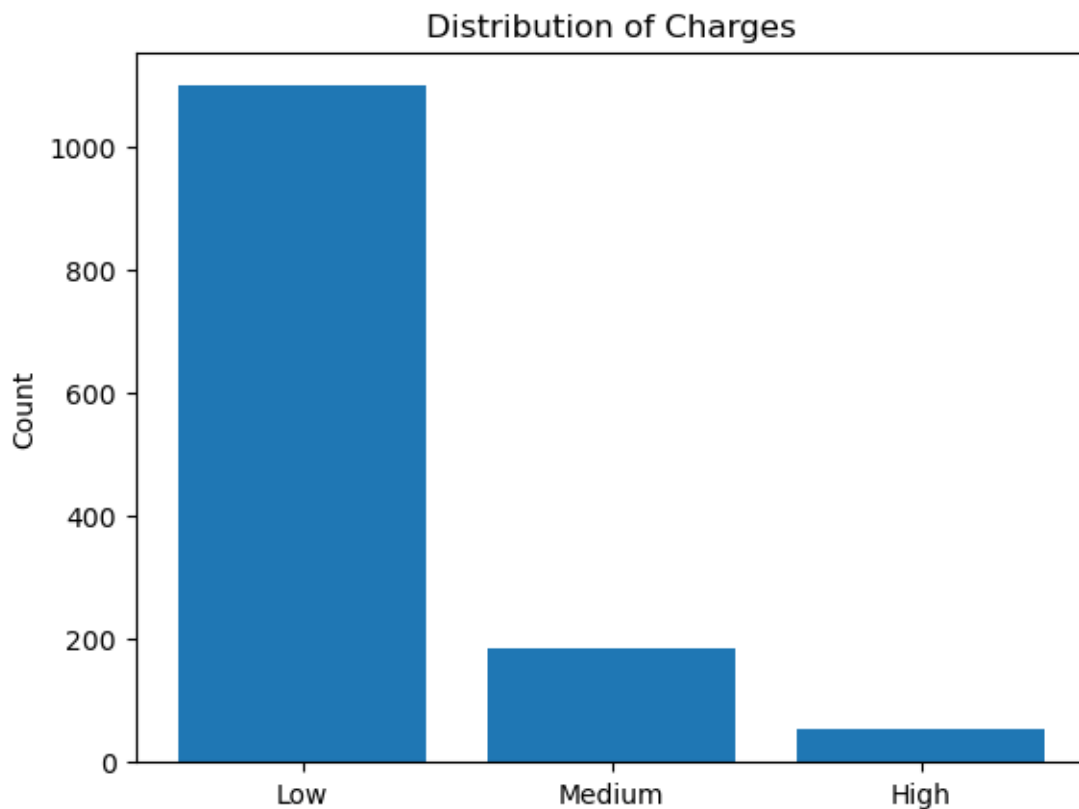
False 1338

Name: count, dtype: int64

III) Creating bins to see the distribution of the charges.

```
[107]: bins = np.linspace(min(df['charges']), max(df['charges']),4)
group_names = ["Low", "Medium", "High"]
df["charges_binned"] = pd.cut(df["charges"], bins, labels = group_names,
    include_lowest= True)
plt.bar(group_names, df["charges_binned"].value_counts())
plt.title("Distribution of Charges")
plt.ylabel("Count")
```

```
[107]: Text(0, 0.5, 'Count')
```



IV) Grouping the dataset by 'sex', 'smoker' and 'region'.

```
[108]: df.groupby('sex').mean(numeric_only=True)
```

```
[108]:
```

	age	bmi	children	charges
sex				
female	39.503021	30.377749	1.074018	12569.578844

```
male      38.917160  30.943129  1.115385  13956.751178
```

```
[109]: df.groupby('smoker').mean(numeric_only=True)
```

```
[109]:
```

	age	bmi	children	charges
smoker				
no	39.385338	30.651795	1.090226	8434.268298
yes	38.514599	30.708449	1.113139	32050.231832

```
[110]: ''' There is a huge disparity between charges for smokers and non-smokers.
↳Smokers on average pay 4 times as much as non-smokers.'''
```

```
[110]: ' There is a huge disparity between charges for smokers and non-smokers. Smokers
on average pay 4 times as much as non-smokers.'
```

```
[111]: df.groupby('region').mean(numeric_only=True)
```

```
[111]:
```

	age	bmi	children	charges
region				
northeast	39.268519	29.173503	1.046296	13406.384516
northwest	39.196923	29.199785	1.147692	12417.575374
southeast	38.939560	33.355989	1.049451	14735.411438
southwest	39.455385	30.596615	1.141538	12346.937377

```
[112]: pearson_coef, p_value = stats.pearsonr(df['bmi'],df["charges"])
print(pearson_coef, p_value)
```

```
0.19834096883362912 2.459085535116604e-13
```

```
[113]: '''There is a statistically significant relationship between the bmi and
↳charges.'''
```

```
[113]: 'There is a statistically significant relationship between the bmi and charges.'
```

V) Feature Engineering.

```
[114]: for column in df.select_dtypes(include='object').columns.tolist():
print(df[column].unique())
```

```
['female' 'male']
['yes' 'no']
['southwest' 'southeast' 'northwest' 'northeast']
```

```
[115]: df = df.drop(columns = "charges_binned")
```

```
[116]: df.head()
```

```
[116]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400

1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
[117]: dataset = pd.get_dummies(data = df,drop_first=True)
```

```
[118]: dataset.head()
```

```
[118]:
```

	age	bmi	children	charges	sex_male	smoker_yes	region_northwest \
0	19	27.900	0	16884.92400	False	True	False
1	18	33.770	1	1725.55230	True	False	False
2	28	33.000	3	4449.46200	True	False	False
3	33	22.705	0	21984.47061	True	False	True
4	32	28.880	0	3866.85520	True	False	True

	region_southeast	region_southwest
0	False	True
1	True	False
2	True	False
3	False	False
4	False	False

```
[119]: dataset.shape
```


```
[119]: (1338, 9)
```

V) Correlation matrix.

```
[120]: dataset_2 = dataset.drop(columns = 'charges')
```

```
[121]: dataset_2.corrwith(dataset['charges']).plot.bar(
    figsize = (16,9),
    title = "Correlation with Charges",
    rot=45)
```


```
[121]: <Axes: title={'center': 'Correlation with Charges'}>
```



```
[122]: corr = dataset.corr()
```

```
[123]: plt.figure(figsize=(20,10))
sns.heatmap(corr, annot= True)
```

```
[123]: <Axes: >
```





VI) Splitting the dataset.

```
[124]: #Independent variable.  
X = dataset.drop(columns= 'charges')
```

```
[125]: X.shape
```

[125]: (1338, 8)

```
[126]: #Dependent variable.  
y = dataset['charges']
```

```
[127]: y.shape
```

[127]: (1338,)

```
[128]: from sklearn.model_selection import train_test_split
```

```
[129]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪ random_state=101)
```

```
[130]: X_train.shape
```

[130]: (1070, 8)

```
[131]: X_test.shape
```

[131]: (268, 8)

VI) Feature Scaling.

```
[132]: from sklearn.preprocessing import StandardScaler
```

```
[133]: scaler = StandardScaler()
```

```
[134]: X_train = scaler.fit_transform(X_train)
       X_test = scaler.transform(X_test)
```

```
[135]: X_train
```

```
[135]: array([[ -1.15786012, -0.68882801, -0.92785237, ..., -0.56943606,
          1.64390454, -0.57087511],
          [-1.51663179,  1.24518419, -0.92785237, ..., -0.56943606,
          1.64390454, -0.57087511],
          [ 0.77950689,  2.31160214, -0.09635988, ..., -0.56943606,
          1.64390454, -0.57087511],
          ...,
          [ 0.77950689,  2.31160214, -0.09635988, ..., -0.56943606,
          1.64390454, -0.57087511],
          [-1.51663179,  1.24518419, -0.92785237, ..., -0.56943606,
          1.64390454, -0.57087511],
          [-1.15786012, -0.68882801, -0.92785237, ..., -0.56943606,
          1.64390454, -0.57087511]])
```

```
[ 1.28178723,  0.5435502 , -0.92785237, ...,  1.75612342,
 -0.60830783, -0.57087511],
 [ 0.92301556,  1.12112479,  0.73513261, ...,  1.75612342,
 -0.60830783, -0.57087511],
 [-0.22505378, -1.78235828, -0.92785237, ..., -0.56943606,
 -0.60830783, -0.57087511]])
```

```
[136]: X_test
```

```
[136]: array([[ -0.08154511,  1.04307417, -0.09635988, ..., -0.56943606,
 -0.60830783, -0.57087511],
 [-1.37312312, -0.31500662, -0.92785237, ..., -0.56943606,
 -0.60830783, -0.57087511],
 [ 1.06652423, -0.86136096, -0.09635988, ..., -0.56943606,
 -0.60830783, -0.57087511],
 ...,
 [-1.15786012, -2.1882215 , -0.09635988, ...,  1.75612342,
 -0.60830783, -0.57087511],
 [ 1.35354156,  0.39402164, -0.92785237, ..., -0.56943606,
 -0.60830783,  1.75169662],
 [ 1.28178723, -0.03402439, -0.92785237, ...,  1.75612342,
 -0.60830783, -0.57087511]])
```

## 5 Part 2: Building the models.

### I) Multiple Linear Regression.

```
[137]: from sklearn.linear_model import LinearRegression
```

```
[138]: model_lr = LinearRegression()
model_lr.fit(X_train, y_train)
```

```
[138]: LinearRegression()
```

```
[139]: y_pred = model_lr.predict(X_test)
```

```
[140]: from sklearn.metrics import r2_score
```

```
[141]: score = r2_score(y_test, y_pred)
```

```
[142]: score
```

```
[142]: 0.760837110132396
```

### II) Random Forest Regressor.

```
[143]: from sklearn.ensemble import RandomForestRegressor
regressor_rf = RandomForestRegressor(random_state=0)
```

```
regressor_rf.fit(X_train, y_train)
```

```
[143]: RandomForestRegressor(random_state=0)
```

```
[144]: y_pred = regressor_rf.predict(X_test)
```

```
[145]: score = r2_score(y_test, y_pred)
```

```
[146]: score
```

```
[146]: 0.8409756227354694
```

III) XB Boost Regressor

```
[147]: from xgboost import XGBRegressor
```

```
[148]: regressor_xgb = XGBRegressor()  
regressor_xgb.fit(X_train, y_train)
```

```
[148]: XGBRegressor(base_score=None, booster=None, callbacks=None,  
                  colsample_bylevel=None, colsample_bynode=None,  
                  colsample_bytree=None, device=None, early_stopping_rounds=None,  
                  enable_categorical=False, eval_metric=None, feature_types=None,  
                  gamma=None, grow_policy=None, importance_type=None,  
                  interaction_constraints=None, learning_rate=None, max_bin=None,  
                  max_cat_threshold=None, max_cat_to_onehot=None,  
                  max_delta_step=None, max_depth=None, max_leaves=None,  
                  min_child_weight=None, missing=nan, monotone_constraints=None,  
                  multi_strategy=None, n_estimators=None, n_jobs=None,  
                  num_parallel_tree=None, random_state=None, ...)
```

```
[149]: y_pred = regressor_xgb.predict(X_test)
```

```
[150]: score = r2_score(y_test, y_pred)
```

```
[151]: score
```

```
[151]: 0.8274000763370111
```

## 6 Part 3: Predict charges for a new customer.

```
[152]: new_customer1_obsv = [[40, 45.5, 4, 1, 1, 0, 0, 0]]
```

```
[153]: regressor_xgb.predict(scaler.transform(new_customer1_obsv))
```

```
[153]: array([44938.145], dtype=float32)
```

```
[154]: new_customer_2 = [[22, 23.5, 4, 1, 1, 0, 0, 0]]
```

```
[155]: regressor_xgb.predict(scaler.transform(new_customer_2))
```

```
[155]: array([16239.731], dtype=float32)
```

```
[ ]:
```